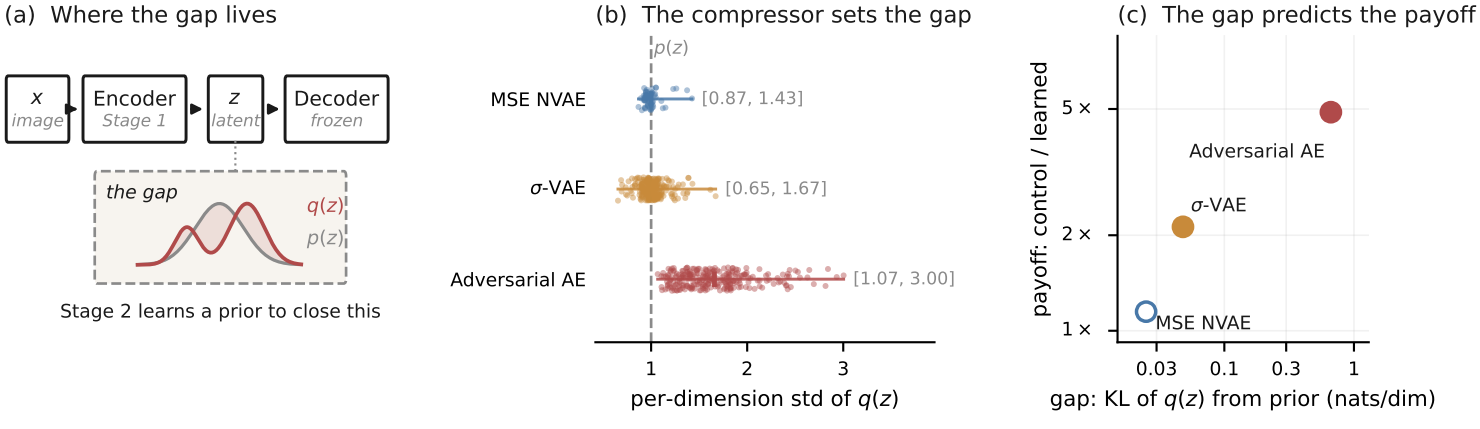
(a) Where the gap lives
x
image
Encoder
Stage 1
z
latent
Decoder
frozen
the gap
q(z)
p(z)
Stage 2 learns a prior to close this
(b) The compressor sets the gap
p(z)
MSE NVAE
[0.87, 1.43]
σ-VAE
[0.65, 1.67]
Adversarial AE
[1.07, 3.00]
1
2
3
per-dimension std of q(z)
(c) The gap predicts the payoff
payoff: control / learned
5 ×
2 ×
1 ×
Adversarial AE
σ-VAE
MSE NVAE
0.03
0.1
0.3
1
gap: KL of q(z) from prior (nats/dim)